

SYSTEMATIC LIQUIDITY RESEARCH

A systematic liquidity-research operation

The deep-tuned book — a per-symbol record across a fifty-two-symbol universe.

ABSTRACT

LIQUIDEX is a systematic research operation built on a single principle: institutional liquidity leaves the same structural footprint wherever it trades. Working from that principle, each symbol in a fifty-two-symbol universe is read down to the institutional behaviour that dominates its order flow and tuned to its own per-symbol recipe through a disciplined optimisation protocol. **Thirty-six are complete** as of this brief: twenty-four clear a profit factor above 1.5, and twelve sit at or below it and are flagged for re-tuning. The strongest cells are long-only fades read through a secular-bull window, so almost every result carries a walk-forward-mandatory caveat. All figures herein are in-sample and net of commission and slippage; their out-of-sample validation is the subject of Section 5.

36 / 52 Symbols deep-tuned	5.75 Peak profit factor (DE)	24 Clear profit factor > 1.5	12 Flagged to re-microtune
--	--	---	--

SECTION 1

The mechanism

Institutional order flow is not noise; it is the residue of decisions. A liquidation cascade, a manufactured stop-run, iceberg absorption defending a level, the exhaustion of a one-sided move — each is a behavioural signature, and each repeats.

The engine identifies the signature, attributes it to the institutional behaviour producing it, and resolves that read into a single rule-based position: entry, a structural stop, a target at the opposing wall, a breakeven and a time-stop. It does not merely detect an anomaly; it profiles the actor behind it — what they are doing, where they are going, and whether they are masking it.

An edge that lives in one market is a curve-fit. An edge that recurs across many is a discovery.

The claim under test is narrow and falsifiable. If the edge were specific to one instrument, it would be an artefact of that instrument's history — an overfit. If instead the same behavioural reads recover an edge across many unrelated instruments, the engine is measuring a structural property of institutional markets rather than the idiosyncrasy of a single chart.

The sections that follow set out, in order: the protocol used to tune each symbol (Section 2); the deep-tuned book across the universe (Section 3); two worked examples of the per-symbol read, with the behavioural reasoning behind each (Section 4); and the validation work that separates these in-sample results from a track record (Section 5).

SECTION 2

The per-symbol protocol

Breadth is meaningless without method. Every symbol is taken through the same protocol. The discipline is deliberately conservative: a behavioural read is admitted only when it improves the risk-adjusted result at its tuned setting, and discarded the moment it does not.

01 Baseline

The engine is reduced to its core — every optional read disabled, default parameters throughout. Raw profit factor, drawdown and trade count are recorded as the benchmark to beat.

02 Greedy forward-selection

Behavioural reads — cascade-stop, opposing-wall target, the directional-risk gate, risk-velocity confirmation, iceberg absorption, liquidity-sweep exhaustion — are added one at a time, retained only if they improve the risk-adjusted result, and reverted otherwise.

03 Coupled parameter tuning

Each retained read is judged at its tuned setting, not its default: its sensitivity is swept in both directions, and a read that cannot pay at any setting is removed rather than left at a default that flatters it.

04 Two-pass refinement

The surviving configuration is walked a second time, each secondary parameter nudged and immediately reverted if the equity curve degrades. The accepted recipe is the one no further change improves.

05 Risk-adjusted scoring

Every decision is taken on $S = PF / (1 + \text{MaxDD})$, never on profit factor alone, subject to a hard floor of thirty trades. A high profit factor on a small sample does not qualify.

06 Out-of-sample gate

A recipe this specific must clear a held-out window before it is treated as real. In-sample strength is the entry condition for that test, not a substitute for it.

We do not search for the setting that makes a chart look good. We search for the read that survives being doubted.

SECTION 3

The deep-tuned book

Each symbol tuned to its own recipe, ranked by profit factor. Thirty-six of fifty-two complete as of **3:33 pm, 20 June 2026**. **PF** ≤ 1.5 is flagged *re-microtune*; $\triangle$ marks walk-forward-mandatory. Full per-symbol configurations are published at wattsonworks.github.io/framework.

TABLE 1 — THE DEEP-TUNED BOOK (36 OF 52 SYMBOLS)

#	SYMBOL	CLASS	PF	MAX DD	WIN	N	ARCHETYPE	STATUS
1	DE $\triangle$	Single-name	5.748	23.6%	87%	76	fade-long	walk-forward
2	GLD $\triangle$	ETF	4.533	18.8%	83%	64	fade-long-ice	walk-forward
3	V $\triangle$	Single-name	4.443	25.4%	81%	80	fade-long	walk-forward
4	COST $\triangle$	Single-name	3.851	23.0%	78%	64	fade-long	walk-forward
5	LRCX $\triangle$	Semis	3.837	26.9%	83%	30	fade-long-cascstop	walk-forward
6	XLV $\triangle$	ETF	3.798	10.5%	87%	46	gated-continuation-fade	walk-forward
7	XLK $\triangle$	ETF	3.184	11.7%	46%	39	fade-long-ice	walk-forward
8	MA $\triangle$	Single-name	2.383	8.9%	36%	33	fade-long-wall-gated	walk-forward
9	QQQ $\triangle$	ETF	2.317	8.7%	47%	32	fade-long-wall	walk-forward
10	ETN $\triangle$	Single-name	2.305	35.2%	77%	60	fade-long	walk-forward
11	DIA $\triangle$	ETF	2.074	10.9%	42%	31	fade-long-wall	walk-forward
12	ASML	Semis	2.054	16.4%	57%	37	fade-both-cascstop-velo	ready
13	INTC $\triangle$	Semis	2.053	46.6%	67%	49	fade-long-cascstop	walk-forward
14	SCHD $\triangle$	ETF	2.047	6.7%	46%	35	fade-long-wall	walk-forward
15	KO	Single-name	1.979	19.8%	54%	201	fade-clean	ready
16	NXPI	Semis	1.971	20.7%	66%	61	fade-gated	ready
17	VOO $\triangle$	ETF	1.849	12.2%	60%	30	fade-clean-cascstop	walk-forward
18	ADI $\triangle$	Semis	1.836	12.7%	42%	31	fade-long-wall-gated	walk-forward
19	KLAC $\triangle$	Semis	1.823	16.5%	41%	51	fade-long-wall	walk-forward
20	MPWR $\triangle$	Semis	1.788	27.2%	55%	31	fade-long-cascstop-velo-bsl	walk-forward
21	PEP	Single-name	1.595	6.5%	49%	107	fade-both	ready
22	TSEM $\triangle$	Semis	1.582	19.2%	44%	39	continuation-long-wall	walk-forward
23	SPY $\triangle$	ETF	1.521	8.5%	38%	32	fade-long-wall	walk-forward
24	AMD $\triangle$	Semis	1.515	21.3%	44%	32	fade-long-wall	walk-forward
25	TXN $\triangle$	Semis	1.455	27.9%	56%	34	fade-long-cascstop-velo-bsl	re-microtune
26	TSM $\triangle$	Semis	1.428	25.2%	40%	30	fade-long-wall	re-microtune
27	ARM $\triangle$	Semis	1.366	31.6%	33%	114	continuation-long-breaker	re-microtune
28	XLI $\triangle$	ETF	1.298	12.5%	40%	50	continuation-long-wall	re-microtune
29	TLT	ETF	1.272	17.7%	49%	205	fade-clean	re-microtune
30	MU $\triangle$	Semis	1.239	13.9%	61%	31	fade-long-velo-nowall	re-microtune
31	NVDA $\triangle$	Semis	1.185	34.0%	33%	222	continuation-long-breaker-gated	re-microtune
32	MRVL	Semis	1.134	43.6%	35%	212	fade	re-microtune
33	QCOM $\triangle$	Semis	1.076	21.9%	24%	38	fade-long-wall-gated	re-microtune
34	MCHP $\triangle$	Semis	1.048	37.2%	27%	84	continuation-long-wall	re-microtune
35	AMAT	Semis	0.944	41.2%	32%	41	fade-long-wall	re-microtune
36	XLFX	ETF	0.895	36.3%	38%	152	fade	re-microtune

wattsonworks.github.io/framework

4

Sixteen of the fifty-two symbols are not yet tuned. The high win rates on the strongest cells are regime artefacts of a secular-bull window — almost every result is walk-forward-mandatory. All figures in-sample, net of commission and slippage.

SECTION 3 (CONTINUED)

What the book shows

Read blind, the book sorts itself into two structural temperaments. **30 of the 36 are faders** — defensives, macro and the broad index complex, where the edge is reverting an exhausted move — and **6 are continuation** names, where it rides the break. The overwhelming majority run **long-only**, the signature of a passive complex accumulating through a rising market. That same fact is the principal caveat: a long-only edge measured through a secular bull is exactly what an out-of-sample window must test.

BY INSTRUMENT CLASS

CLASS	SYMBOLS	READY	MEAN PF	BEST
Index & sector ETFs	11	8	2.25	4.53 · GLD
Semiconductors	18	9	1.63	3.84 · LRCX
Single-name leaders	7	7	3.19	5.75 · DE

The queue, and what discipline holds back

Twelve names sit at or below a profit factor of 1.5 and are **held back for re-tuning** rather than presented as finished. Two — AMAT and the financials proxy XLF — could not clear a profit factor of 1.0 at thirty trades and are deferred outright. No result is admitted on fewer than thirty trades, and no figure has been selected after the fact to flatter the whole. Sixteen of the fifty-two symbols remain untuned; the book is published as it stands, not as it will read when it is finished.

SECTION 4

Worked examples: the per-symbol read

Two symbols from the book, opened all the way up. Together they make the same point from opposite sides: the read is not a setting but an understanding of the actor's temperament, expressed as a configuration.

PEP — *when the desk is patient*

A liquidation cascade never comes for a consumer staple. PepsiCo is not where leverage blows up; it is where patient, mandate-driven capital accumulates and defends. Its flow is quiet, two-sided and mean-reverting — and that is precisely why a default read fails it, listening for violence in a market that trades on patience.

Reading the desk meant matching the algorithm to the psychology: both directions enabled, the reward-to-risk gate raised so only the cleanest reversals qualify, iceberg absorption set deliberately slow to wait for genuinely defended size, and the risk gate tightened on the short side. The coin-flip became a clean, low-drawdown edge.

PEP • RESULT

Profit factor	1.595
Net return	+31.24%
Max drawdown	6.47%
Win rate	48.60%
Trades	107

Configuration: long + short, raised R:R gate, iceberg at 5 hits, high-volume sweep exhaustion, short-side risk gate tightened. Verified live.

TLT — *when the move never reaches the wall*

TLT is long-duration Treasuries — a macro vehicle, not a tape that liquidity desks push to clean technical walls. Its price is dragged by rate expectations and central-bank language, and the full move to the opposing wall rarely completes, because a data print interrupts it first. Everything one adds to it is fragile: tighten the stop and the chop removes you; widen the target and the move dies before it arrives.

What rescued it was profit-taking discipline — banking 60% of the position two-thirds of the way to the wall — and a wide sweep stop. A buffer-30 stop lifts the result past the earlier 1.232; the whole edge balances on those two levers.

TLT • RESULT

Profit factor	1.272
Net return	+18.05%
Max drawdown	17.66%
Win rate	48.78%
Trades	205

The levers: a wide sweep stop (buffer 30) over the earlier 1.232, plus Partial-TP banking 60% at 0.66 to the wall. Long + short, cascade-fade, wall-target. Verified live. **Flagged to re-microtune** (PF ≤ 1.5).

A recipe this specific is where overfitting hides. Both cells carry the out-of-sample flag of Section 5: an in-sample result is the entry condition for validation, never a deployment claim.

SECTION 5

Limitations and forward work

The figures in this document are in-sample. The programme below is the standard by which the research is held — the work that separates a strong backtest from a track record. The same discipline is applied to the research itself.

THE VALIDATION STANDARD

TEST	STATUS	NOTE
Commission & slippage	Modelled	Applied to every backtest; all figures are net of costs, not gross.
Capacity	Structural	The book is liquid large-caps and major ETFs that absorb size; the precise ceiling is not yet formally measured.
Out-of-sample	Pending — next	A held-out window the engine has never seen; the immediate priority, and the gate every deep-tuned recipe waits behind.
Walk-forward	Planned	Rolling re-optimisation and re-test, to show each per-symbol recipe is stable through time rather than fitted to one window.
Monte Carlo	Planned	Trade-order and return resampling, to bound the distribution of outcomes and the risk of ruin under sizing.
Live record	The bar	A forward, time-stamped track record — the only dispositive evidence for a fund or licensing conversation.

The strongest cells in the book are long-only fades through a secular-bull window; their high win rates are the clearest sign that out-of-sample work is owed, not optional. The honest summary of the present state is this: the evidence is strong enough to justify that work, and not yet a substitute for it. Twelve of the thirty-six are explicitly held back for re-tuning; sixteen of the fifty-two are not yet begun.

SECTION 6

The system, and access

The research is delivered as three instruments, invite-only on TradingView. One engine decides the trade; two companions make its ranking and its risk legible.

LIQUIDEX
THE ENGINE

5,000+ lines of Pine v6. Reads liquidation cascades, order flow, order blocks, iceberg walls and liquidity sweeps, and resolves them into one position — entry, structural stop, and a target at the opposing wall.

Keeper
LIVE RANKING

Runs the engine across the whole book and ranks every name live by profit factor, momentum and risk — so the operator sees the leader the engine prefers, not a guess.

Prophet
RISK GATE

Makes directional risk visible — a composite score, safe and danger zones, and a live position-size verdict that trims exposure as risk rises and stands the operator down in the danger zone.

ACCESS & PARTNERSHIP

Operators receive invite-only access on TradingView: send a TradingView username, access is granted, and the signals arrive as alerts on phone, desktop or a connected trading bot. Access and pricing are available on request.

The defensible asset is the method — a protocol for reading the institutional desks behind the order flow that reproduces across the universe, equities and macro alike. It is open to licensing, a strategic research partnership, or full acquisition.

WhatsApp +972 50 678 7586 · wattsonworks.github.io/framework